


Problem Session 4

Problem 4-1. Sequence Rotations

Below is a Sequence AVL Tree T . Perform operation $T.delete_at(8)$ and draw the tree after each rotation operation performed during the operation.

(SEE NEXT PAGE)



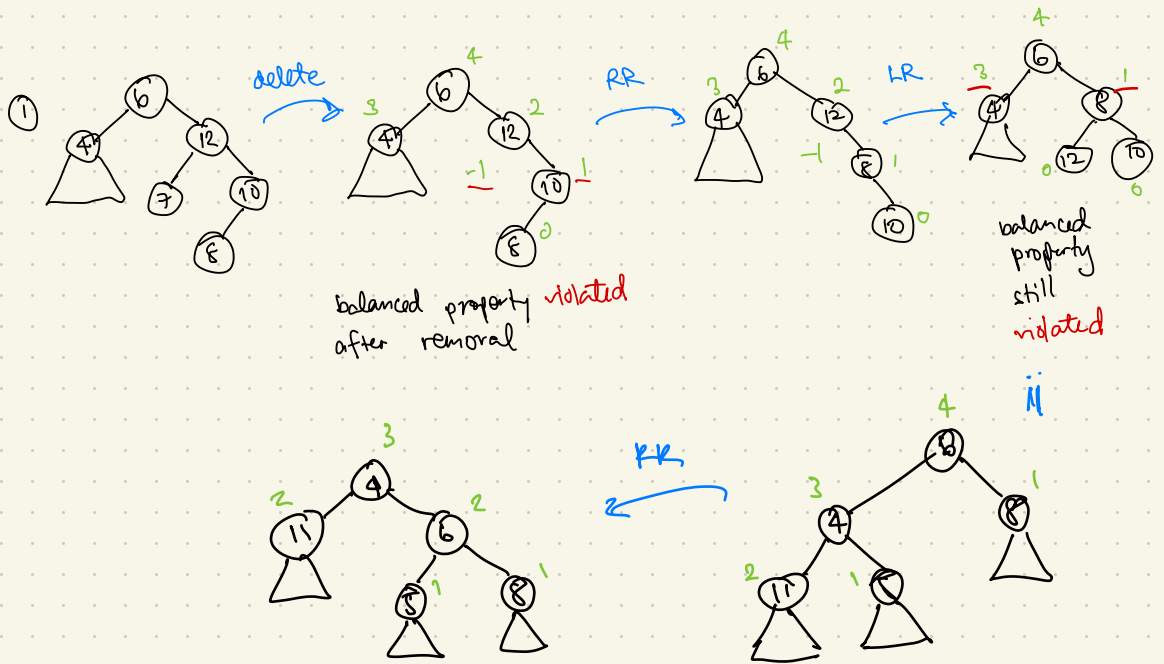
Problem 4-2. Fick Nury

Fick Nury directs an elite group of n superheroes called the Revengers. He has heard that supervillain Sanos is making trouble on a distant planet and needs to decide whether to confront her. Fick surveys the Revengers and compiles a list of n polls, where each poll is a tuple matching a different Revenger's name with their integer opinion on the topic. Opinion $+s$ means they are for confronting Sanos with strength s , while opinion $-s$ means they are against confronting Sanos with strength s . Fick wants to generate a list containing the names of the $\log n$ Revengers having the strongest opinions (breaking ties arbitrarily), so he can meet with them to discuss. For this problem, assume that the record containing the polls is **read-only** access controlled (the material is classified), so any computation must be written to alternative memory.

(a) Describe an $O(n)$ -time algorithm to generate Fick's list.

(SEE NEXT PAGE)

(b) Now suppose Fick's computer is only allowed to write to at most $O(\log n)$ space. Describe an $O(n \log \log n)$ -time algorithm to generate Fick's list.



② We can "build" a max-heap in-place by treating the polls in an array as already have "inserted" each poll. (if need be we can copy to new $\Theta(n)$ memory in $\Theta(n)$ time.) Then we can heapify down (starting at the leaves) each node, resulting in $\Theta(n)$ build time of a max-heap. Note, the keys are tuples $(i, |s_i|)$ for each Reverger i . Once the max-heap priority queue is built, we can repeatedly pop $\log n$ items, which takes $O(\log^2 n)$ time, which is still asymptotically faster than $O(n)$, thus, the overall runtime is $O(n)$ in the worst case.

2 Generally correct for traditional auction mechanics wherein bidders cannot decrease their bids, but they can here!

Problem Session 4

Problem 4-3. SCLR

Stormen, Ceiserson, Livest, and Rein are four academics who wrote a very popular textbook in computer science, affectionately known as SCLR. They just found k first editions in their offices, and want to auction them off online for charity. Each bidder in the auction has a unique integer bidder ID and can bid some positive integer amount for a single copy (but may increase or decrease their bid while the auction is live). Describe a database supporting the following operations, assuming n is the number of bidders in the database at the time of the operation. For each operation, state whether your running time is worst-case, expected, and/or amortized.

<code>new_bid(d, b)</code>	record a new bidder ID d with bid b in $O(\log n)$ time
<code>update_bid(d, b)</code>	update the bid of existing bidder ID d to bid b in $O(\log n)$ time
<code>get_revenue()</code>	return revenue from selling to the current k highest bidders in $O(1)$ time

Problem 4-4. Receiver Roster

Coach Bell E. Check is trying to figure out which of her football receivers to put in her starting lineup. In each game, Coach Bell wants to start the receivers who have the highest **performance** (the average number of points made in the games they have played), but has been having trouble because her data is incomplete, though interns do often add or correct data from old and new games. Each receiver is identified with a unique positive integer jersey number, and each game is identified with a unique integer time. Describe a database supporting the following operations, each in **worst-case** $O(\log n)$ time, where n is the number of games in the database at the time of the operation. Assume that n is always larger than the number of receivers on the team.

④ <code>record(g, r, p)</code>	record p points for jersey r in game g
① <code>clear(g, r)</code>	remove any record that jersey r played in game g
② <code>ranked_receiver(k)</code>	return the jersey with the k^{th} highest performance

STORE POINTERS TO EACH OF

BELOW IN HASH TABLE OR

ARRAY or something for fast lookups

① Each jersey r has set AVL tree keyed on g , with tree augmented by size, total points for games they've played in. Calculating performance is $O(1)$ at root

② Need some sort of sequence AVL, where insertion is done by traversing D.S. keyed by performance

3

- insert on d in $O(\log n)$
- update (d, b) in $O(\log n)$
- get k-largest b in $O(1)$

could just use
hash table
 $d \rightarrow b$

maintain separate
APL tree on bid
amounts with size, sum
augmentation, constrained
to $\leq k$ nodes

on each
insert and
update?

Problem 4-5. Warming Weather

Gal Ore is a scientist who studies climate. As part of her research, she often needs to query the maximum temperature the earth has observed within a particular date range in history, based on a growing set of measurements which she collects and adds to frequently. Assume that temperatures and dates are integers representing values at some consistent resolution. Help Gal evaluate such range queries efficiently by implementing a database supporting the following operations.

<code>record_temp(t, d)</code>	record a measurement of temperature t on date d
<code>max_in_range(d1, d2)</code>	return max temperature observed between dates d_1 and d_2 inclusive

To solve this problem, we will store temperature measurements in an AVL tree with binary search tree semantics keyed by date, where each node A stores a measurement $A.item$ with a date property $A.item.key$ and temperature property $A.item.temp$.

- (a) To help evaluate the desired range query, we will augment each node with:

$A.max_temp$, the maximum temperature stored in A 's subtree; and both $A.min_date$ and $A.max_date$, the minimum and maximum dates stored in A 's subtree respectively. Describe a $O(1)$ -time algorithm to compute the value of these augmentations on node A , assuming all other nodes in A 's subtree have already been correctly augmented.

- (b) A subtree **partially overlaps** an inclusive date range if the subtree contains at least one measurement that is within the range **and** at least one measurement that is outside the range. Given an inclusive date range, prove that for any binary search tree containing measurements keyed by dates, there is at most one node in the tree whose left and right subtrees both partially overlap the range.

- (c) Let `subtree_max_in_range(A, d1, d2)` be the maximum temperature of any measurement stored in node A 's subtree with a date between d_1 and d_2 inclusive (returning `None` if no measurements exist in the range). Assuming the tree has been augmented as in part (a), describe a **recursive** algorithm to compute the value of `subtree_max_in_range(A, d1, d2)`. If h is the height of A 's subtree, your algorithm should run in $O(h)$ time when A partially overlaps the range, and in $O(1)$ time otherwise.

- (d) Describe a database to implement operations `record_temp(t, d)` and `max_in_range(d1, d2)`, each in **worst-case** $O(\log n)$ time, where n is the number of unique dates of measurements stored in the database at the time of the operation.

- (e) Implement your database in the Python class `Temperature_DB` extending the `Set_AVL_Tree` class provided; you will only need to implement parts (a) and (c) from above.

(skipped for time)

SEE
NEXT
FEW
PAGES

Right idea,
poorly articulated
argument!

⑤ A For computing 'max-temp' property in $O(1)$ time, we can simply take the max of the 'max-temp' property from 'A's' temperature, the left child's 'max-temp' property, and the right child's 'max-temp' property (if they exist). Similarly, 'A.min-date' can be computed in $O(1)$ time by taking its left child's 'min-date' property (if it exists) otherwise defaulting to 'A.item.key' (its own date). Lastly, $O(1)$ computation of 'max-date' for is the same except with the right child's 'max-date' or 'A.item.key' if A has no right child.

⑥ Proof: Notice that any interval that has at least one end either $\leq \text{min. key}$ or $\geq \text{max key}$ will have 0 nodes satisfying $P(n) := \text{node } n \text{'s}$ left and right subtrees are partially overlapped. The same is true for intervals where both ends are out of range of the keys spanned by the BST. We can restrict ourselves to cases where the entire interval is spanned by the key range of the BST. We can proceed by establishing a contradiction if we assume $\exists n. P(n) \Rightarrow \exists m \neq n. P(m)$.

For a node n to satisfy $P(n)$ an interval must start \leq n 's left child's key AND end \geq n 's right child's key. The start must also be \geq the min keyed node in n 's left subtree AND The end $<$ the max keyed node in n 's right subtree.

Thus, one interval that results in satisfying $P(n)$ for node n cannot satisfy $P(n^*)$ for a different node n^* because two subtrees are rooted by exactly one node (by definition). But this contradicts $\exists m \neq n \ P(m)$.

⑤ ② We can immediately return a value for cases when 1) the date interval is out of range of the Set AVL or 2) the interval spans all dates in the Set AVL. For the former we can return 'None' in $O(1)$ time and for the latter, return 'A.maxTemp' in $O(1)$ time, too. For the partial overlap case, we can achieve $O(h)$ performance:

(after checking base cases above)

if $A.item.date > d_2$: recurse on A.left

if $A.item.date < d_1$: recurse on A.right

else: $\#$ A in the interval $[d_1, d_2]$

if $d_1 \leq A.min.date$: return the max of
 $A.item.temp, A.left.max.temp$, recurse on A.right

if $d_2 \geq A.max.date$: return the max of
 $A.item.temp, A.right.max.temp$, recurse on A.left

else: $\#$ A's left/right subtrees partially overlap $[d_1, d_2]$
 return max of recurse on A.left,
 recurse on A.right, A.item.temp

one child's subtree entirely contained in $[d_1, d_2]$

at most, there is one node that suffices, so at most 2 paths down from this node

This runs in $O(h)$ time since $O(1)$ work at each touched node!

\Leftarrow

forget to keep
max if data
already in database!

⑤ We can continue to use a Set AVL tree as described in the problem. 'record-temp' is basically an 'insert' on a set AVL, taking care to perform any rotations needed to preserve the AVL property and also updating the augmented metadata of all ancestors. Rotations are $O(1)$ time operations each and so are subtree augmentation updates. These need to be done $O(\log n)$ in the worst case (this is a height-balanced tree, so $h = O(\log n)$). 'max-in-range' is simply calling 'subarray-max-in-range' (from 1c) on the root of this database. We know if AVL property is maintained on all writes, this runs in $O(h) = O(\log n)$ worst-case time.