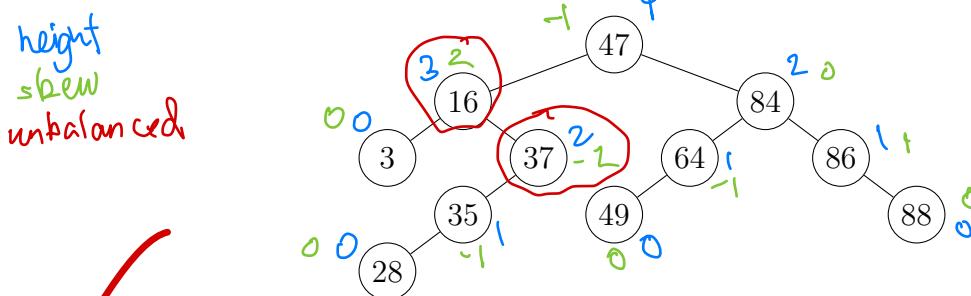


Problem Set 4

Please write your solutions in the L^AT_EX and Python templates provided. Aim for concise solutions; convoluted and obtuse descriptions might receive low marks, even when they are correct.

Problem 4-1. [10 points] Binary Tree Practice

- ✓ (a) [2 points] The Set Binary Tree T below is **not height-balanced** but does satisfy the **binary search tree** property, assuming the key of each integer item is itself. Indicate the keys of all nodes that are not height-balanced and compute their skew.

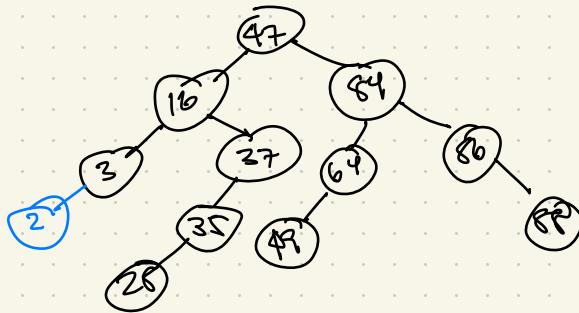


- ✓ (b) [5 points] Perform the following insertions and deletions, one after another in sequence on T , by adding or removing a leaf while maintaining the binary search tree property (a key may need to be swapped down into a leaf). For this part, **do not** use rotations to balance the tree. Draw the modified tree after each operation.

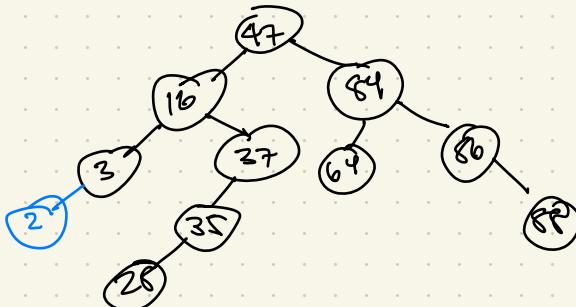
```
1 T.insert(2)
2 T.delete(49)
3 T.delete(35)
4 T.insert(85)
5 T.delete(84)
```

- ✓ (c) [3 points] For each unbalanced node identified in part (a), draw the two trees that result from rotating the node in the **original** tree left and right (when possible). For each tree drawn, specify whether it is height-balanced, i.e., all nodes satisfy the AVL property.

1 1B

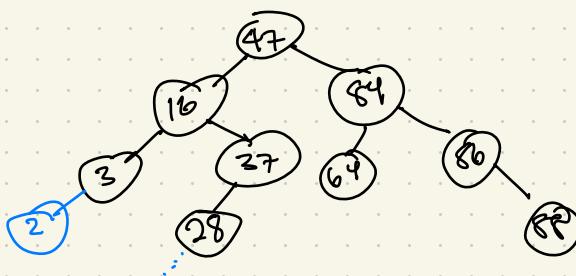


T.insert(2)
(simple insert since
③.predecessor is empty)



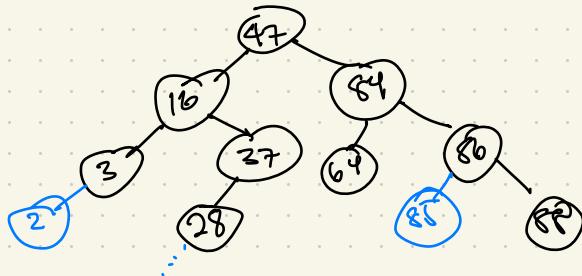
T.delete(49)

(simple removal since
it's a leaf)



T.delete(35)

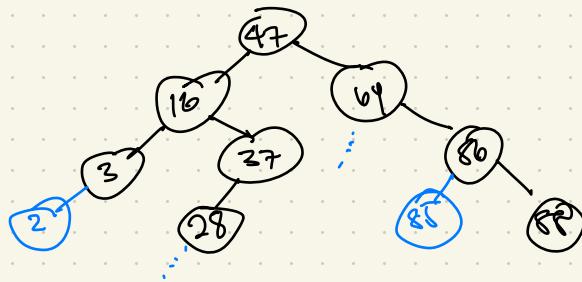
(switch key with predecessor
key, delete leaf)



T.insert(85)

(insert before 86) \Leftrightarrow
insert after 84, which
calls 84.successor/
84.right_subtree.first
which returns 86 w/ no
left child, so insert there)

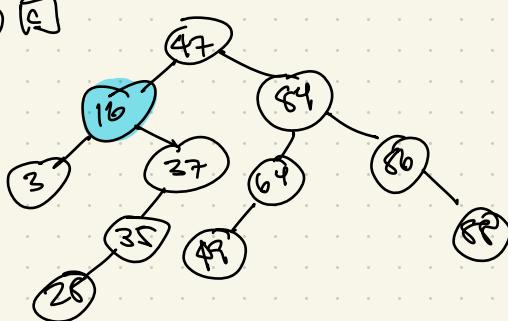
(①) (CONT.)



7. delete (84)

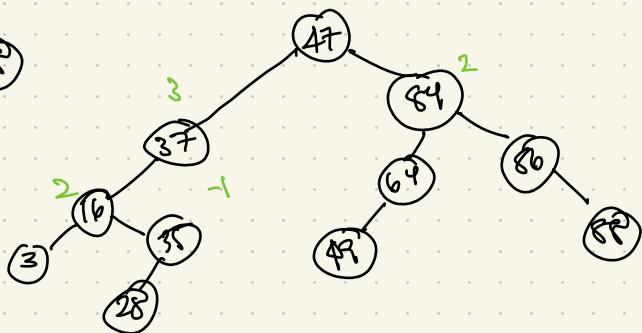
switch keys 86 (84) and
81, predecessor() since 84 left
exists; then simply detach
the leaf wtr key 84

① (a)



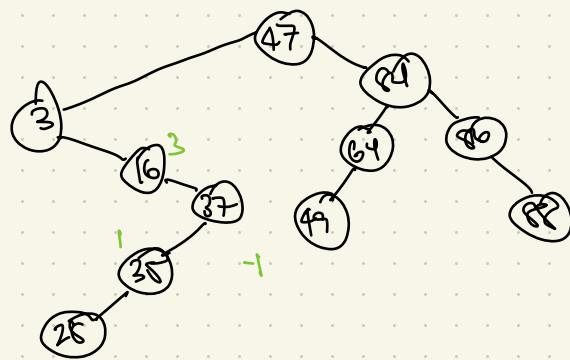
LR
16

(height)



RR
16

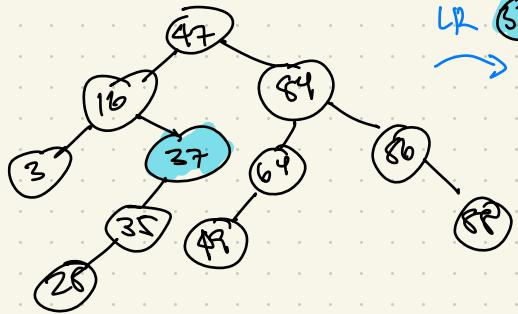
not height-balanced
after LR 16



not height-balanced
after RR 16

(1c) cont.)

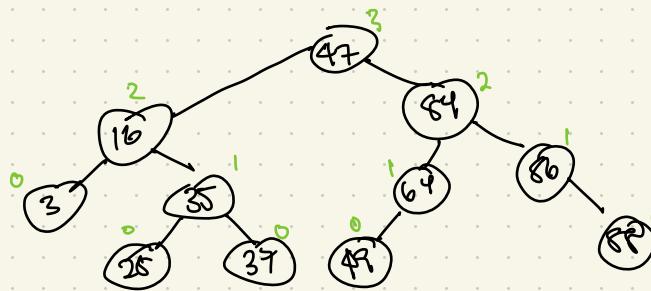
(height)



LR (37)

cannot lf (37) b/c
(37) has no right child
subtree!

LR 27



height-balanced after LR 27 !

Note: Material on this page requires material that will be covered in **L08 on March 3, 2020**. We suggest waiting to solve these problems until after that lecture. All other pages of this assignment can be solved using only material from L07 and earlier.

Problem 4-2. Heap Practice [10 points]

For each array below, draw it as a **complete**¹ binary tree and state whether the tree is a max-heap, a min-heap, or neither. If the tree is neither, turn the tree into a min-heap by repeatedly swapping items that are **adjacent in the tree**. Communicate your swaps by drawing a sequence of trees, marking on each tree the pair that was swapped.

- (a) [4, 12, 8, 21, 14, 9, 17]
- (b) [701, 253, 24, 229, 17, 22]
- (c) [2, 9, 13, 8, 0, 2]
- (d) [1, 3, 6, 5, 4, 9, 7]

Problem 4-3. [10 points] Gardening Contest

Gardening company Wonder-Grow sponsors a nation-wide gardening contest each year where they rate gardens around the country with a positive integer² **score**. A garden is designated by a **garden pair** (s_i, r_i) , where s_i is the garden's assigned score and r_i is the garden's unique positive integer **registration number**.

- (a) [5 points] To support inclusion and reduce competition, Wonder-Grow wants to award identical trophies to the top k gardens. Given an unsorted array A of garden pairs and a positive integer $k \leq |A|$, describe an $O(|A| + k \log |A|)$ -time algorithm to return the registration numbers of k gardens in A with highest scores, breaking ties arbitrarily.
- (b) [5 points] Wonder-Grow decides to be more objective and award a trophy to every garden receiving a score strictly greater than a reference score x . Given a max-heap A of garden pairs, describe an $O(n_x)$ -time algorithm to return the registration numbers of all gardens with score larger than x , where n_x is the number of gardens returned.

I did BFS, but solution presents DFS
 (both work in $O(n_x)$ time).

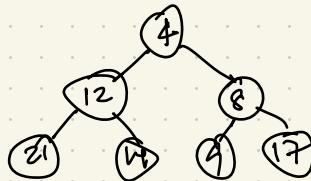
¹Recall from Lecture 8 that a binary tree is **complete** if it has exactly 2^i nodes of depth i for all i except possibly the largest, and at the largest depth, all nodes are as far left as possible.

²In this class, when an integer or string appears in an input, without listing an explicit bound on its size, you should assume that it is provided inside a constant number of machine words in the input.

2

A

[4, 12, 8, 21, 14, 9, 17]

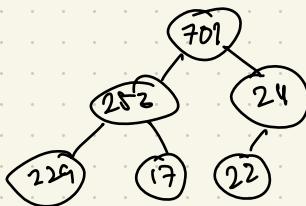


min-heap

B

[701, 253, 24, 229, 17, 22]

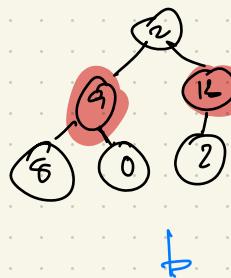
max-heap



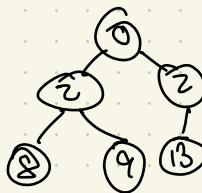
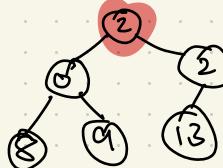
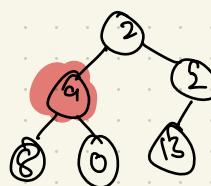
C

[2, 9, 13, 8, 0, 2]

To heapify in-place, better to min-heapify-down from bottom up to get $O(n)$ build time!



min-heap property violated!



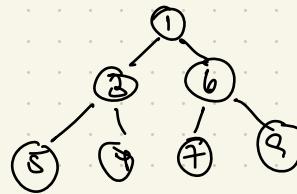
min-heap

[0, 2, 2, 8, 9, 12]

②

[1, 3, 6, 5, 4, 9, 7]

min-heap



③ ~~A~~ Build a max-heap of the gardens using the score as key in $O(|A|)$ time. Then pop + delete k gardens from the heap in $O(k \log |A|)$ time.

~~B~~ We can traverse the heap from the root using breadth-first search: maintain a queue of nodes whose children we have yet to check for the property that their key $\geq x$. Start with root, check for desired property. If yes, add root to result array, and add child nodes to said queue. If no, there is no need to process the rest of the heap (due to heap property invariant). Keep repeating this process until the queue is empty and return the result array. This will run in $O(nx)$ time (since using dynamic arrays for queue and result array).

Problem 4-4. [15 points] Solar Supply

Entrepreneur Bonty Murns owns a set S of n solar farms in the town of Fallmeadow. Each solar farm $(s_i, c_i) \in S$ is designated by a unique positive integer **address** s_i and a farm **capacity** c_i : a positive integer corresponding to the maximum energy production rate the farm can support. Many buildings in Fallmeadow want power. A building (b_j, d_j) is designated by a unique **name** string b_j and a **demand** d_j : a positive integer corresponding to the building's energy consumption rate.

To receive power, a building in Fallmeadow must be connected to a **single** solar farm under the restriction that, for any solar farm s_i , the sum of demand from all the buildings connected to s_i may not exceed the farm's capacity c_i . Describe a database supporting the following operations, and for each operation, specify whether your running time is worst-case, expected, and/or amortized.

| | |
|--|--|
| <code>initialize(S)</code> | Initialize database with a list $S = ((s_0, c_0), \dots, (s_{n-1}, c_{n-1}))$ corresponding to n solar farms in $O(n)$ time. |
| <code>power_on(b_j, d_j)</code> | Connect a building with name b_j and demand d_j to any solar farm having available capacity at least d_j in $O(\log n)$ time (or return that no such solar farm exists). |
| <code>power_off(b_j)</code> | Remove power from the building with name b_j in $O(\log n)$ time. |
| <code>customers(s_i)</code> | Return the names of all buildings supplied by the farm at address s_i in $O(k)$ time, where k is the number of building names returned. |

lazy in
the analysis

Problem 4-5. [15 points] Robot Wrangling

Dr. Squid has built a robotic arm from $n+1$ rigid bars called **links**, each connected to the one before it with a rotating joint (n joints in total). Following standard practice in robotics³, the orientation of each link is specified locally relative to the orientation of the previous link. In mathematical notation, the change in orientation at a joint can be specified using a 4×4 **transformation matrix**. Let $\mathcal{M} = (M_0, \dots, M_{n-1})$ be an array of transformation matrices associated with the arm, where matrix M_k is the change in orientation at joint k , between links k and $k+1$.

To compute the position of the **end effector**⁴, Dr. Squid will need the arm's **full transformation**: the ordered matrix product of the arm's transformation matrices, $\prod_{k=0}^{n-1} M_k = M_0 \cdot M_1 \cdot \dots \cdot M_{n-1}$. Assume Dr. Squid has a function `matrix_multiply(M_1, M_2)` that returns the matrix product⁵ $M_1 \times M_2$ of any two 4×4 transformation matrices in $O(1)$ time. While tinkering with the arm changing one joint at a time, Dr. Squid will need to re-evaluate this matrix product quickly. Describe a database to support the following **worst-case** operations to accelerate Dr. Squid's workflow:

| | |
|---|--|
| <code>initialize(\mathcal{M})</code> | Initialize from an initial input configuration \mathcal{M} in $O(n)$ time. |
| <code>update_joint(k, M)</code> | Replace joint k 's matrix M_k with matrix M in $O(\log n)$ time. |
| <code>full_transformation()</code> | Return the arm's current full transformation in $O(1)$ time. |

³More on forward kinematic robotics computation here: https://en.wikipedia.org/wiki/Forward_kinematics

⁴i.e., the device at the end of a robotic arm: https://en.wikipedia.org/wiki/Robot_end_effector

⁵Recall, matrix multiplication is not commutative, i.e., $M_1 \cdot M_2 \neq M_2 \cdot M_1$, except in very special circumstances.

④ We will rely on 3 data structures:

1. set AVL whose nodes are keyed on available capacity and that contain a pointer to a hash table of the attached buildings \mapsto demand
2. dictionary mapping farm names to their respective node in ①
3. dictionary mapping building to node of farm they are attached to.

Since the capacities are positive integers we can actually sort them in linear time (using radix sort), then construct ① in $O(n)$ time by storing the middle item as the root and then recursively building the left + right subtrees from the left + right fans in the sorted sequence. Thus, `initialize(s)` takes $O(n)$ time overall and the tree produced has height $O(\log n) \Rightarrow$ balanced \Leftrightarrow AVL property. ② results in $O(n)$ time. `power-on(b_j, d_j)` can be implemented in $O(\log n)$ time by finding a node in ① with sufficient capacity in $O(\log n)$, adding itself to that node's attached building hash set in $O(1)$ time, reducing that node's available capacity and then deleting and

(④ cont.) reinserting the node into the correct place in $2 \cdot O(\log n) + O(1) = O(\log n)$ time and updating B_j in $O(1)$ time.

'power-off(b_j)' requires looking b_j up in B ($O(1)$ time), going to its associated farm node and removing the building from the attached buildings hash table ($O(1)$ time), adding d_j to available capacity ($O(1)$ time), and deleting / re-inserting this farm's node in $O(\log n)$ time overall.

'customers(s_i)' requires a lookup to farm i 's node in R ($O(1)$ time), then returning all b_i in that node's attached building hash table. Thus, $O(k)$ time overall for this operation.

⑤ We can accomplish all required operations under the given constraints by building a height-balanced binary tree that's ~~self~~ tree augmented such that every node carries information about the matrix product of its left child, itself, and right child (recursively, for all nodes in the tree). In particular, 'initialize(M)' can build, in $O(n)$

(R) cont.)

time, the desired tree from the sequence of matrices M (by assigning M_{112} to the root, and recursing on the left & right halves of M accordingly). This preserves the necessary order for the matrix products. The subtree augmentation can also happen in this initialization. $\text{update_point}(k, M)$ effectively performs binary search on the tree (as long as we store a "size" property at each node) to find M_k . Updating the matrix product property for all of node k 's ancestors is $O(\log n)$, so overall we achieve $O(\log n)$ time for this operation. Finally, the subtree augmentation pays off for $O(1)$ time 'full-transformation': we just return the matrix product at the root node. (Note that all runtimes described above are worst-case.)

Problem 4-6. [40 points] $\pi z^2 a$ Optimization

Liza Pover has found a Monominos pizza left over from some big-TEX recruiting event. The pizza is a disc⁶ with radius z , having n **toppings** labeled $0, \dots, n - 1$. Assume z fits in a single machine word, so integer arithmetic on $O(1)$ such integers can be done in $O(1)$ time. Each topping i :

- is located at Cartesian coordinates (x_i, y_i) where x_i, y_i are integers from range $R = \{-z, \dots, z\}$ (you may assume that **all coordinates are distinct**), and
- has integer **tastiness** $t_i \in R$ (note, topping tastiness can be negative, e.g., if it's pineapple⁷).

Liza wants to pick a point (x', y') and make a pair of cuts from that point, one going straight down and one going straight left, and take the resulting *slice*, i.e., the intersection of the pizza with the two half-planes $x \leq x'$ and $y \leq y'$. The tastiness of this slice is the sum of all t_i such that $x_i \leq x'$ and $y_i \leq y'$. Liza wants to find a **tastiest** slice, that is, a slice of maximum tastiness. Assume there exists a slice with **positive tastiness**.

- ~~(a)~~ (a) [2 points] If point (x', y') results in a slice with tastiness $t \neq 0$, show there exists $i, j \in \{0, 1, \dots, n - 1\}$ such that point (x_i, y_j) results in a slice of equal tastiness t (i.e., a tastiest slice exists resulting from a point that is both vertically and horizontally aligned with toppings).

- ~~(b)~~ (b) [8 points] To make finding a tastiest slice easier, show how to modify a Set AVL Tree so that:

- it stores **key-value items**, where each item x contains a value $x.val$ (in addition to its key $x.key$ on which the Set AVL is ordered);
- it supports a new tree-level operation `max_prefix()` which returns in **worst-case** $O(1)$ time a pair $(k^*, \text{prefix}(k^*))$, where k^* is any key stored in the tree T that maximizes the **prefix sum**, $\text{prefix}(k) = \{x.val \mid x \in T \text{ and } x.key \leq k\}$ (that is, the sum of all values of items whose keys are $\leq k$); and
- all other Set AVL Tree operations maintain their running times.

- ~~(c)~~ (c) [5 points] Using the data structure from part (b) as a black box, describe a **worst-case** $O(n \log n)$ -time algorithm to return a triple (x, y, t) , where point (x, y) corresponds to a slice of maximum tastiness t .

- (d) [25 points] Write a Python function `tastiest_slice(toppings)` that implements your algorithm from part (c), including an implementation of your data structure from part (b).

There is a bug in the solution code; see trivial example in test case I wrote (last one) in the code!

⁶The pizza has thickness a , so it has volume $\pi z^2 a$.

⁷If you believe that Liza's Pizza preferences are objectively wrong, feel free to assert your opinions on Piazza.

Basic idea is that the sort must be done on y values for toppings with shared x !

⑥ A slice given by (x', y') with tastiness $\neq 0$ will have at least one topping with nonzero tastiness. i, j can be found by taking the "top" and "right" - most toppings with tastiness $\neq 0$. (Note, this might be the same topping.) Given the definition of a slice, the slice given by (x_i, y_j) has the same tastiness as the slice given by (x', y') since any toppings not found in slice (x_i, y_j) will have zero tastiness and all toppings in (x', y') are also in (x_i, y_j) .

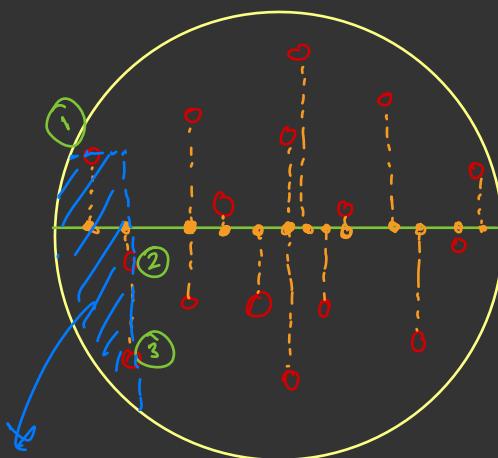
⑥ Argue the subtree with sum & prefix-max: such that for any node, prefix max is given by:

$$\max(A.\text{left_prefix_max},$$

$$A.\text{left_sum} + A.\text{val}$$

$$A.\text{left_sum} + A.\text{val} + A.\text{right_prefix_max})$$

⑥ We can order all toppings by x , iterate through this ordering constructing set $A[k]$ described above keeping on y and keeping track of sum, prefix sum on tastiness t . $O(n \log n)$ to sort, $O(n \log n)$ to construct this data structure. Note, we can simply hold onto the x^*, y^* on each iteration. Getting t^* is $O(1)$ per above. (All runtimes are worst case.)



This "slice", when inserted into the set AVL tree according to the given order, would consider slice w/ ① & ② as a potential tastiest slice!



For toppings with shared x -coordinates, they also must be inserted in y -coordinate order for 'tastiest_slice' function to work as written in the solution!

```

1  from Set_AVL_Tree import BST_Node, Set_AVL_Tree
2
3  class Key_Val_Item:
4      def __init__(self, key, val):
5          self.key = key
6          self.val = val
7
8      def __str__(self):
9          return "%s,%s" % (self.key, self.val)
10
11 class Part_B_Node(BST_Node):
12     def subtree_update(A):
13         super().subtree_update()
14         ######
15         # ADD ANY NEW SUBTREE AUGMENTATION HERE #
16         #####
17
18 class Part_B_Tree(Set_AVL_Tree):
19     def __init__(self):
20         super().__init__(Part_B_Node)
21
22     def max_prefix(self): # O(i) worst case
23         """
24             Output: (k, s) | a key k stored in tree whose
25                         | prefix sum s is maximum
26         """
27         k, s = 0, 0
28         #####
29         # YOUR CODE HERE #
30         #####
31         return (k, s)
32
33 def tastiest_slice(toppings): # O(n log n) worst case
34     """
35         Input: toppings | List of integer tuples (x,y,t) representing
36                         | a topping at (x,y) with tastiness t
37         Output: tastiest | Tuple (X,Y,T) representing a tastiest slice
38                         | at (X,Y) with tastiness T
39     """
40     B = Part_B_Tree()    # use data structure from part (b)
41     X, Y, T = 0, 0, 0
42     #####
43     # YOUR CODE HERE #
44     #####
45     return (X, Y, T)

```

MIT OpenCourseWare

<https://ocw.mit.edu>

6.006 Introduction to Algorithms

Spring 2020

For information about citing these materials or our Terms of Use, visit: <https://ocw.mit.edu/terms>